

Controlling the cloud with Python

Pycon.it
9 maggio 2009

Luca Mearelli

<http://luca.im>

@lmea



kiaraservice.com



*-as-a-service
elastic & fluid
scalable
programmable



EC2

S3

CloudFront

SimpleDB

SQS

Elastic MapReduce

MechanicalTurk

...



<http://code.google.com/p/boto>

```
python setup.py install
export AWS_ACCESS_KEY_ID=<your key>
export AWS_SECRET_ACCESS_KEY=<your secret>
# ~/.boto
# /etc/boto.cfg
```



S3 - Simple Storage Service

storage service

scalable

replicated & distributed

web aware (HTTP / bittorrent)

flexible security

eventually consistent



S3 - Connecting

```
import boto
conn = boto.connect_s3()
conn = S3Connection()

conn.is_secure
conn = boto.connect_s3(is_secure=False)
```



S3 - Buckets

```
bucket = conn.create_bucket('pycon3')  
rs = conn.get_all_buckets()  
for bucket in rs :  
    print bucket.name  
  
bucket.delete()  
conn.delete_bucket('pycon3')
```



S3 - Keys & Objects

```
from boto.s3 import Key
key = Key(bucket)
key = bucket.new_key('aaaa')

key.name = 'aaaa'
key.exists()
```



S3 - Keys & Objects

```
key.set_contents_from_string(  
    'My Test',  
    headers={'Content-Type': 'text/plain'})
```

```
key.set_contents_from_filename('test.txt')
```



S3 - Keys & Objects

```
rs = bucket.get_all_keys()  
rs = bucket.get_all_keys(prefix='test/level1/',  
                           delimiter='/',  
                           maxkeys=5)  
key = bucket.get_key('demo/foo')
```



S3 - Keys & Objects

```
key = bucket.get_key('foo')
```

```
fp = open('foo.txt', 'wb')
```

```
key.get_contents_to_file(fp)
```

```
s = key.get_contents_as_string()
```



S3 - Keys & Objects

```
key.set_metadata('format', 'avi')  
key.get_metadata('format')
```



S3 - Permissions

```
bucket.make_public(recursive=True)
key.make_public()
key.generate_url(10)
key.set_acl('public-read-write')
CannedACLStrings = ['private',
                    'public-read',
                    'public-read-write',
                    'authenticated-read']
```



S3 - Logging

```
lb = conn.create_bucket('pycon3.logging')  
pb = conn.create_bucket('pycon3.public')  
lb.set_as_logging_target()  
pb.enable_logging(lb)
```



Computing on demand

Any application on any OS

Persistent/ephemeral storage

Server customization

Flexible security



EC2 - Connecting

```
conn = boto.connect_ec2()  
conn = EC2Connection()  
  
regions = boto.regions()  
conn = regions[0].connect()
```



```
rs = conn.get_all_images(owners=['012345678901'])
```

```
image = conn.get_image(image_id='ami-5647a33f')
```



EC2 - Keypairs

```
key_pair = conn.create_key_pair('my_key')  
key_pair.save('/Users/luca/.ssh')
```



EC2 - Instances

```
res = conn.run_instances('ami-5647a33f',  
                          instance_type='m1.small',  
                          key_name='my_key',  
                          min_count=1, max_count=1,  
                          security_groups=['web'],  
                          user_data=None)  
  
inst = res.instances[0]
```



EC2 - Instance

```
while not instance.update() == 'running':  
    print instance.state  
    time.sleep(5)  
  
print inst.id  
print inst.public_dns_name  
print inst.private_dns_name  
print inst.state  
print inst.key_name  
print inst.launch_time
```



EC2 - Instance

```
inst.stop()  
res.stop_all()
```

```
inst.reboot()
```

```
console = inst.get_console_output()  
print console.instance_id  
print console.timestamp  
print console.output
```




```
group = conn.create_security_group(  
    group_name='web',  
    group_desc='Web Servers')  
  
rs = conn.get_all_security_groups()
```



```
conn.authorize_security_group(  
    'web',  
    ip_protocol='tcp',  
    from_port='80',  
    to_port='80',  
    cidr_ip='0.0.0.0/0']
```

```
group.revoke('tcp', 80, 80, '0.0.0.0/0']
```



EC2 - Elastic block storage

```
ebs = conn.create_volume(size,  
                          zone,  
                          snapshot=None)  
  
conn.attach_volume(volume_id=ebs.volume_id,  
                   instance_id=inst.id,  
                   '/sdb')  
  
conn.create_snapshot(volume_id)
```



Distributed queue

Web scale

Redundant infrastructure



SQS - Connecting

```
conn = boto.connect_sqs()  
conn = SOSConnection()
```



SQS - Queue

```
queue = conn.create_queue('myqueue')
queue.url
# 'https://queue.amazonaws.com/591131556747/myqueue'
queue.get_timeout()
queue.set_timeout(120)
rs = conn.get_all_queues()
queue = conn.get_queue('myqueue')
```



SQS - Messages

```
from boto.sqs import Message

msg = Message()
msg.set_body('A test message')
msg.attributes['timestamp'] = 1202131

queue.write(msg)
```



SQS - Messages

`Message()`

`MHMessage()`

`RawMessage()`

`JSONMessage()`

`queue.set_message_class(MHMessage)`



SQS - Messages

```
queue.count()
```

```
rs = queue.get_messages(num_messages=1,  
                        visibility_timeout=None)
```

```
msg = queue.read(visibility_timeout=None)
```

```
msg.get_body()
```



SQS - Messages

```
msg.delete()
```

```
queue.clear()
```



SQS - Dumping & loading

```
queue.save_to_file(fp, sep='\n')  
queue.load_from_file(fp)
```

```
queue.save_to_s3(bucket, prefix=None)  
queue.load_from_s3(bucket, prefix=None)
```



Structured data storage

Schema-free

Queryable



SDB - Connecting

```
conn = boto.connect_sdb()  
conn = SDBConnection()
```



SDB - Domains

```
dom = conn.create_domain('pycon3')
```

```
dom = conn.lookup('pycon3')
```

```
rs = conn.get_all_domains()
```



SDB - Domains

```
md = conn.get_metadata()  
md.item_count
```



SDB - Items

```
item = dom.new_item('item1')
item['k1'] = 'value'
item['k2'] = 10
item['k3'] = ['a', 'b', 'c']
item.add_value('k3', 'd')
item.save()
```



SDB - Read & query

```
dom.get_item['item1']
```

```
rs = dom.query("[ 'k1' = 'value' ]")
```

```
rs = dom.select(  
    "select * from pycon3 where k1 = 'value'")
```

```
for item in rs :  
    print item.name
```



SDB - Query examples

```
"['city' = 'Seattle' or 'city' = 'Portland']"
```

```
"['author' starts-with 'Robert']"
```

```
"['author' does-not-start-with 'Robert']"
```

```
"['first name' = 'John'] intersection ['last name' = 'Smith']"
```

```
"['tag' starts-with 'Amazon'] union ['description' = 'SimpleDB']"
```

```
"not ['country' = 'USA' or 'country' = 'UK']"
```



SDB - Select examples

```
"select * from mydomain where city = 'Seattle' or  
city = 'Portland'"
```

```
"select * from mydomain where author like 'Rob%'"
```

```
"select * from mydomain where year is not null"
```

```
"select * from mydomain where every[keyword] in  
['Book', 'Paperback'] "
```

```
"select itemName[] from mydomain"
```

```
"select count[*] from mydomain"
```



SDB - Dumping to xml

```
d = dom.to_xml()
```



???



@lmea #pycontre #aws

Grazie!

Luca :-)



@lmea #pycontre #aws

[http://lmea-docs.s3.amazonaws.com/
controlling_the_cloud_with_python](http://lmea-docs.s3.amazonaws.com/controlling_the_cloud_with_python)



@lmea #pycontre #aws